

Model Name (Single)	Data Count	Training Time (s)	Avg Inference Time	Overall	Response Adaptation	Context Mistakes (/30)	Failures (/30)	Mistakes (/30)	Notes
System prompt	-	-	3.01	4.5	3.5	0	0	2	Long sentences, no variation. Narration
Embeddings	70	-	2.91	4.5	4.5	1	0	0	Bryn thinks he is John
Embeddings	150	-	3.12	4	3.5	2	1	1	Longer responses. Bryn thinks he is John + that he knows Jenny
Embeddings	1500	-	2.41	4	4	1	1	0	Longer sentences but more natural.
Finetune - Instruct - QLoRA	30	8.5943	3.87	3	3	8	0	2	Very long sentences, narration. Mistakes on village name, neighbors
Finetune - Instruct - QLoRA	70	16.6934	3.99	3	3	5	0	7	Lots of narration.
Finetune - Instruct - QLoRA	150	9.0604	3.75	3	3.5	7	0	6	Lots of narration.
Finetune - Instruct - LoRA	30	8.6111	1.19	4.5	4	5	1	0	Failures on < im_start >assistant. Slightly longer responses.
Finetune - Instruct - LoRA	150	38.3922	1.03	4.5	4.5	3	0	0	Context mistakes on neighbors.
Finetune - Base - QLoRA	1500	313.1691	1.6	3.5	4.5	5	1	0	Context mistakes on all questions about neighbors
Finetune - Base - LoRA	1500	309.3864	0.58	4	4.5	2	0	0	Context mistakes only on direct questions (Who is Jenny?)